


Heart Disease Science

Finn B

2/9/2021



Introduction

The purpose of this report is the analysis and methodology of several health data of patients from 1988. The data shows if a patient has **heart disease**. It describes a range of conditions that are in connection with the heart. The data set used is a data set provided by **Donald Bren School of Information and Computer Sciences** from the **University of California, Irvine** originally. This project will concentrate on a database from the **V.A. Medical Center, Long Beach and Cleveland Clinic Foundation** created by **Robert Detrano, M.D., Ph.D.**. The data was sourced from **Kaggle**, where the data was initially processed.

Origin of this database: [Archive.ics.uci](https://archive.ics.uci.edu/)

First step is exploration and cleaning of the dataset. After that, each attribute is examined for its relation to the target variable. In the section of modeling, several classification methods are run through to predict whether a patient has heart disease or not. The methods used are **logistic regression, decision tree, random forest, support vector machines** and **k-nearest neighbors**.

Data exploration and cleaning

Data exploration

This report excludes 62 attributes from the original database to work only with a subset of 14 attributes, containing **13 features** and **one outcome variable** to consider if a patient has heart disease. The database contains health data of **303 patients**.

On a first view you can see what features will accompany the final outcome variable in this project. Before heading into the analysis we need to understand what the different attributes tell us:

age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
57	0	0	120	354	0	1	163	1	0.6	2	0	2	1
57	1	0	140	192	0	1	148	0	0.4	1	0	1	1

Data cleaning

For data cleaning, some data from the original data base was changed. The levels of 'sex' were changed to 'female' and 'male'. The levels of 'target' were changed to 'disease' and 'no disease' to have a quiet better overview. Furthermore some of the attributes were encoded as factors to enable a better work: **sex**, **cp**, **fbs**, **restecg**, **exang**, **slope**, **ca**, **thal**, **disease(target)**.

```
HeartData <- HeartData %>%
  mutate(sex=ifelse(sex==0, 'female', 'male'))
HeartData$sex <- as.factor(HeartData$sex)
```

```
HeartData$cp <- as.factor(HeartData$cp)
HeartData$cp <- revalue(HeartData$cp, c('0'='asymptomatic',
                                         '1'='atypical angina',
                                         '2'='non anginal pain',
                                         '3'='typical angina'))
```

```
HeartData <- HeartData %>%
  mutate(fbs=ifelse(fbs==1,
                    '>120',
                    '<=120'))
HeartData$fbs <- as.factor(HeartData$fbs)
```

```
HeartData$restecg <- as.factor(HeartData$restecg)
HeartData$restecg <- revalue(HeartData$restecg, c('0'='left ventricular hypertrophy',
                                                    '1'='normal',
                                                    '2'='st-t abnormality'))
```

```
HeartData <- HeartData %>%
  mutate(exang=ifelse(exang==0,
                      'no',
                      'yes'))
HeartData$exang <- as.factor(HeartData$exang)
```

```
HeartData$slope <- as.factor(HeartData$slope)
HeartData$slope <- revalue(HeartData$slope, c('0'='downsloping',
                                             '1'='flat',
                                             '2'='upsloping'))
```

```
HeartData$ca <- as.factor(HeartData$ca)
HeartData$ca <- revalue(HeartData$ca, c('4'=NA))
```

```
HeartData$thal <- as.factor(HeartData$thal)
HeartData$thal <- revalue(HeartData$thal, c('0'=NA,
                                             '1'='fixed defect',
                                             '2'='normal',
                                             '3'='reversible defect'))
```

```
HeartData <- HeartData %>%
  mutate(target=ifelse(target==0,
                        "disease",
                        "no disease"))
HeartData$target <- as.factor(HeartData$target)
HeartData$disease <- HeartData$target
HeartData$target <- NULL
attr(HeartData, 'spec') <- NULL
```

Attribute	Meaning
age	Patients age (29-77 years)
sex	Female (0) and Male (1)
cp - chest pain type	asymptomatic (0); atypical angina (1); non-anginal pain (2); typical angina (3)
trestbps - resting blood pressure	in mm/Hg on admission to the hospital ¹
chol - serum cholesterol	in mg/dl
fbs - fasting blood sugar	> 120 mg/dl; no(0) yes(1)
restecg - resting electrocardiographic results	probable or definite left ventricular hypertrophy by Estes' criteria(0); normal(1); having ST-T wave abnormality(2)
thalach	maximum heart rate achieved
exang - exercise induced angina	no(0); yes(1)
oldpeak	ST depression induced by exercise relative to rest
slope - slope of peak exercise ST segment	downsloping(0); flat(1); upsloping(2)
ca - number of major vessels colored by flourosopy	vessels(0-3); NA(4)
thal - Thallium Stress Test Result	NA(0); fixed defect(1); normal(2); reversible defect(3)
disease - angiographic disease status	> 50% diameter narrowing (0); < 50% diameter narrowing (1)

¹ Judging from the values, the **systolic pressure** (the pressure when the heart pushes blood out) is given here.

age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	disease
63	male	typical angina	145	233	>120	left ventricular hypertrophy	150	no	2.3	downsloping	0	fixed defect	no disease
37	male	non anginal pain	130	250	<=120	normal	187	no	3.5	downsloping	0	normal	no disease
41	female	atypical angina	130	204	<=120	left ventricular hypertrophy	172	no	1.4	upsloping	0	normal	no disease
56	male	atypical angina	120	236	<=120	normal	178	no	0.8	upsloping	0	normal	no disease
57	female	asymptomatic	120	354	<=120	normal	163	yes	0.6	upsloping	0	normal	no disease
57	male	asymptomatic	140	192	<=120	normal	148	no	0.4	flat	0	fixed defect	no disease

Data analysis

In this part of the project we will dig deeper into the attributes and potential effects on the disease. But first we will have a look on the categorization of disease and on the most obvious and superficial indicators: Age and Sex.

Disease

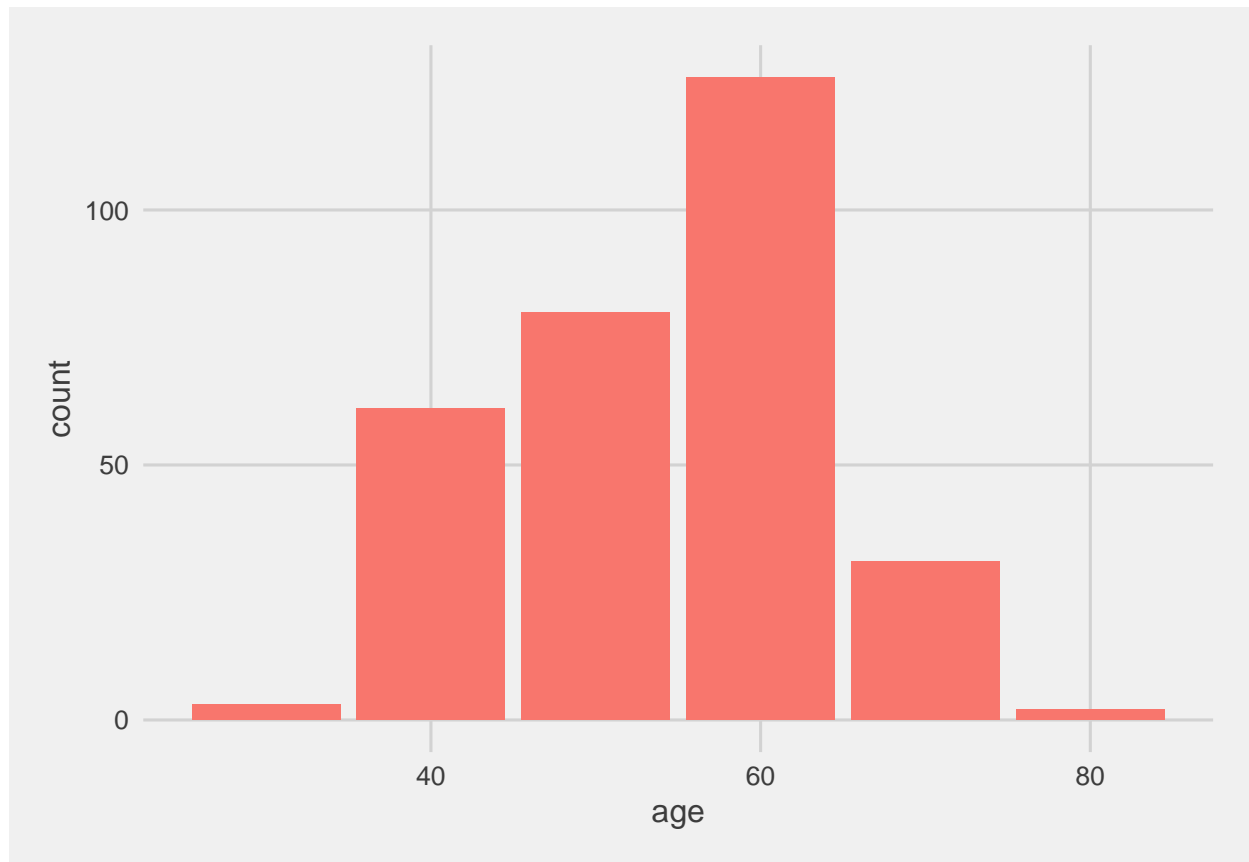
To determine the diagnosis of heart disease status the **angiographic disease** status was used. This was differentiated into two conditions that differ in the percentage diameter narrowing of <50% and >50% of coronary arteries.

disease	cases
disease	138
no disease	165

As we can see the proportion of patients with disease also called prevalence was at 45.5% in the database.


Age

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	29.0	47.5	55.0	54.4	61.0	77.0



The age range goes from 29 years to 77 year. The median age is at 55 years, while we can see that the most patients are between 55 and 65 years.

This distribution can also be determined by the distribution of patients divided into disease and no disease where most patients with disease are in this range of age.



Sex

sex	count	disease
female	96	0.250
male	207	0.551

The distribution by sex is dominated by male patients, around 68% of the patients are male. The mean age of female patients is quite lower than the mean age of male patients.


This can be seen in the mean of patients **with** heart diseases as well.

female:


##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	43.0	56.8	60.5	59.0	62.0	66.0

male:

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	35.0	51.0	57.5	56.1	61.0	77.0




As we can see in the following plot there are much more female patients without heart disease than with heart disease. Furthermore the number of male patients with and without diseases seem to be similar while there are more male patients around 60 with heart disease.



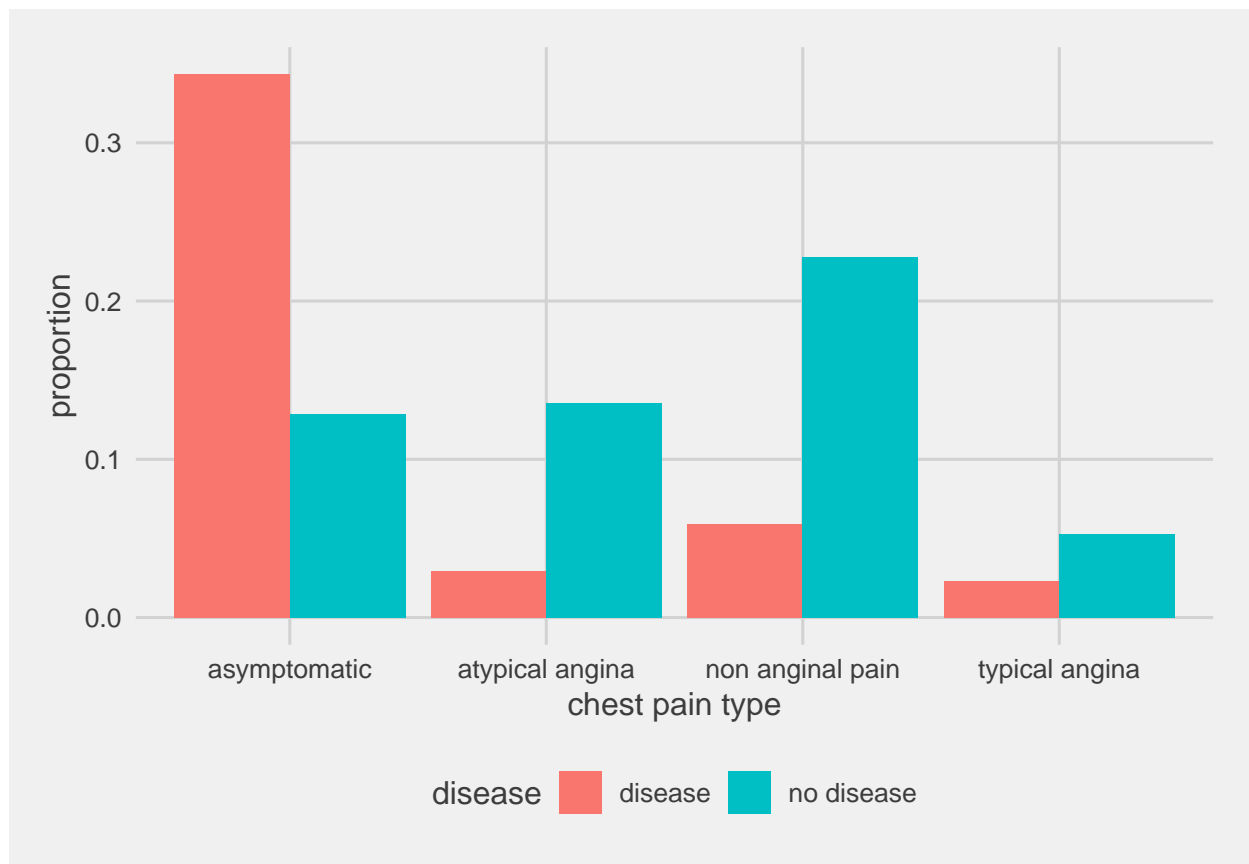
Chest pain type

As previously researched, there are four types of chest pain. **Asymptomatic** pain means that a patient has no symptoms/pain. **Angina** is a pain that the patient has near the heart, often described as chest tightness. Angina pain is categorized into **atypical** angina and **typical** angina. Most patients data shows asymptomatic and non anginal pain. Only a small amount of patients had typical angina:



Something that may not have been expected by many is the following result. Except of asymptomatic pain the proportion of patients with disease is far below 50% for each category of pain.

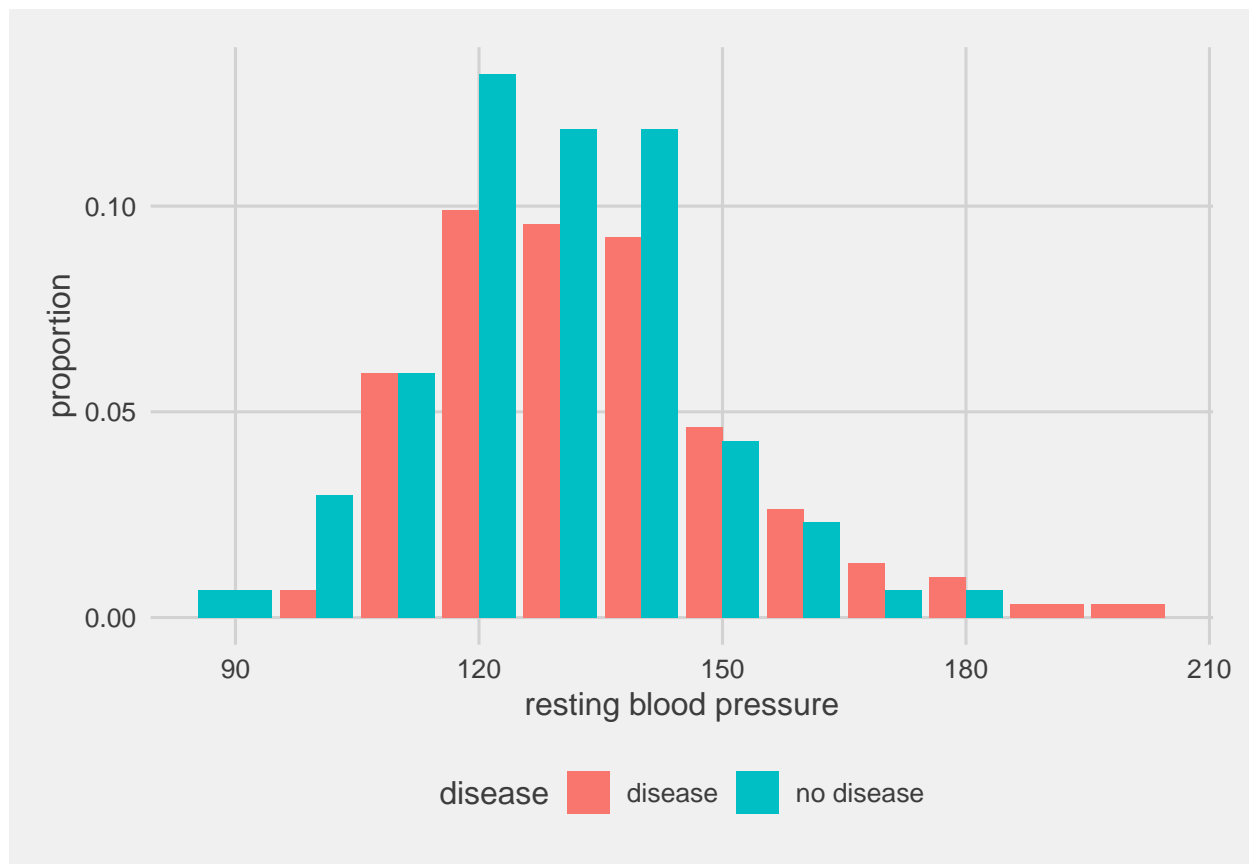
chest pain	disease (prop)
asymptomatic	0.727
atypical angina	0.180
non anginal pain	0.207
typical angina	0.304




Resting blood pressure

For most patients (68%) the blood pressure is higher than the **ideal systolic blood pressure** of between 90 and 120mm/Hg. In the second plot we can observe a slightly increasing proportion of disease with a higher systolic resting blood pressure.

trestbps <= 120mm/Hg	trestbps > 120mm/Hg
0.32	0.68

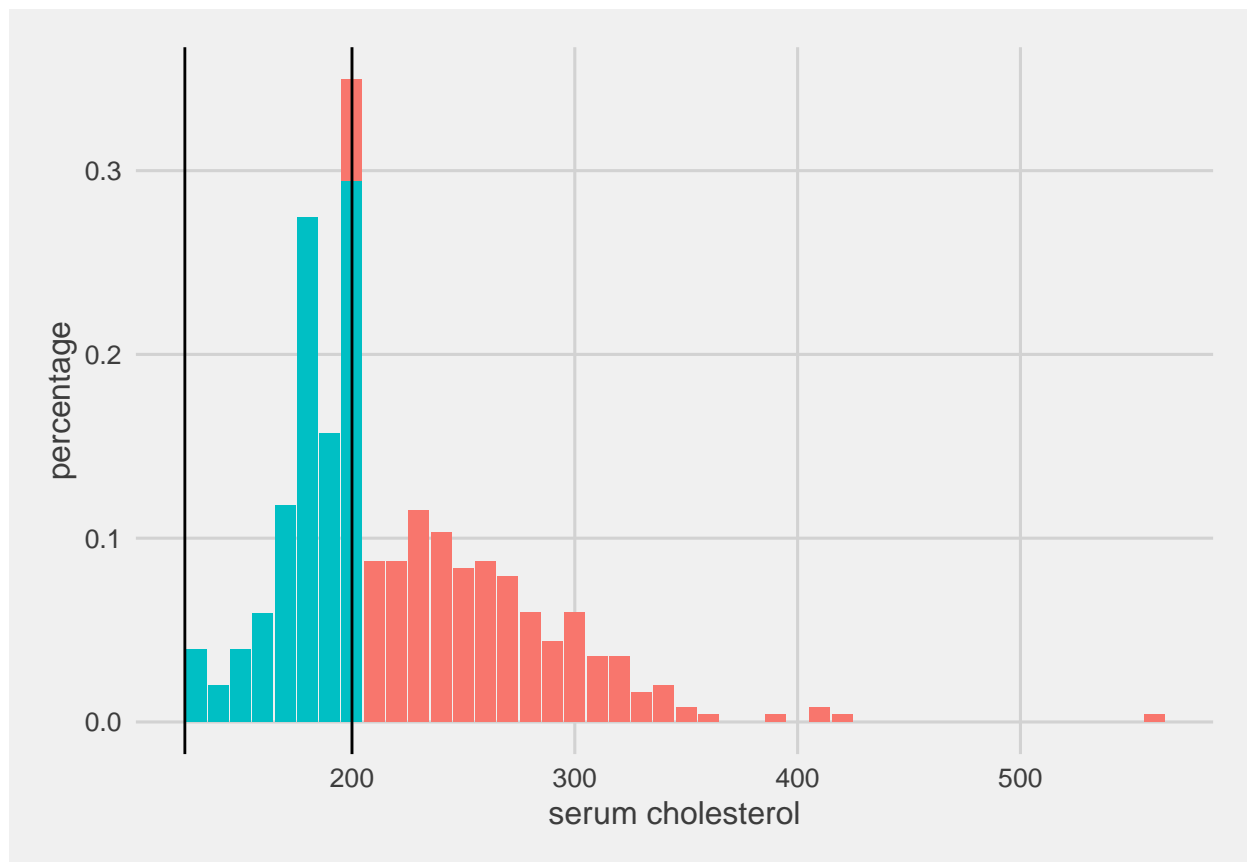




Serum cholesterol

Since there is too little information about what type of cholesterol level is given we assume total cholesterol. **Healthy cholesterol level** for adults is between 125mg/dL and 200mg/dL.

Most patients serum cholesterol is higher than the healthy range:



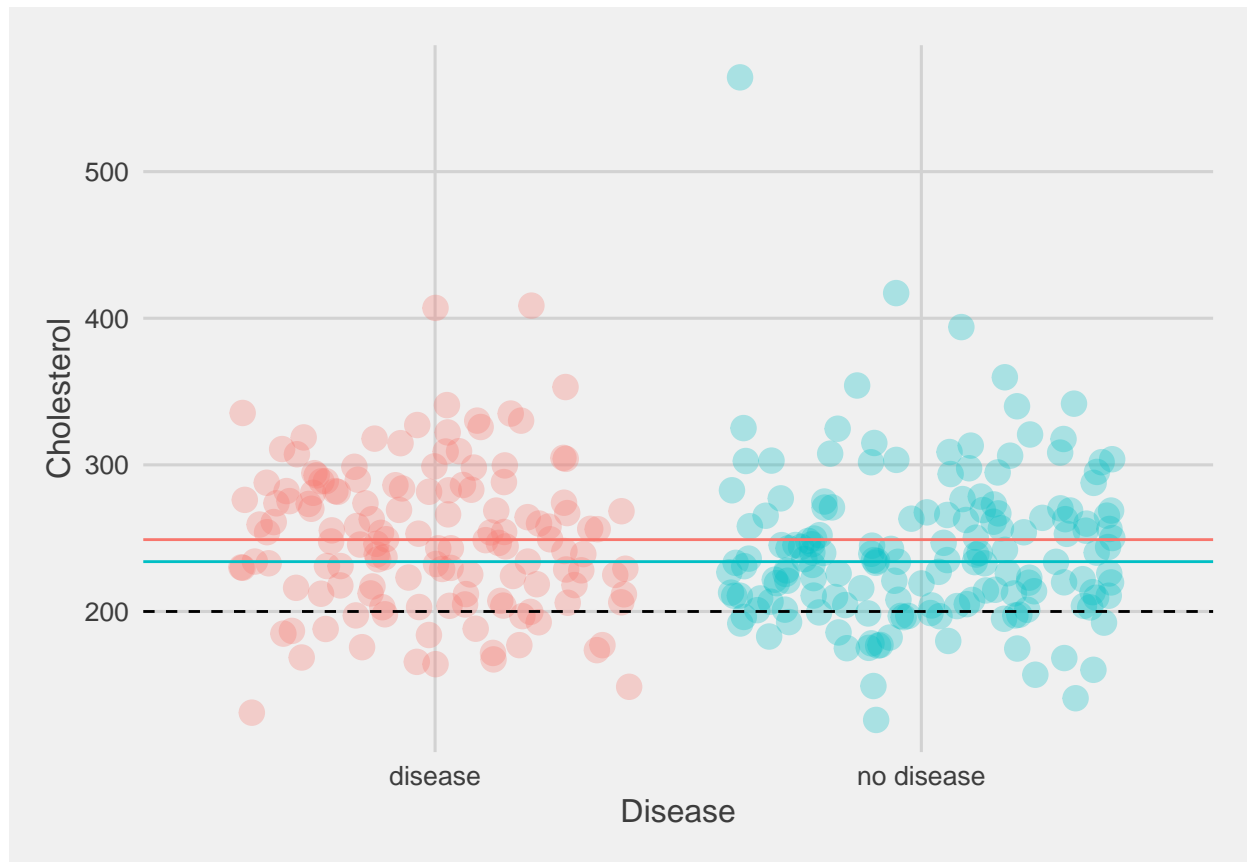
This can also be seen in comparison between diseased and healthy patients. The median of cholesterol of both groups is higher than 200 while cholesterol in the group of diseased patients is slightly higher.

```
HD.chol.median.mean <- HeartData %>%
  group_by(disease) %>%
  summarize(me=mean(chol), med=median(chol))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
HD.chol <- HeartData %>%
  group_by(disease) %>%
  select(disease, chol)

ggplot(data=HD.chol, aes(disease, chol, color=disease)) +
  geom_jitter(width = 0.4, alpha = 0.3, size=4) +
  stat_smooth(method="lm", formula=disease~1, se=FALSE) +
  geom_hline(data=HD.chol.median.mean, aes(yintercept = med, color=disease)) +
  geom_hline(yintercept=200, linetype = "dashed") +
  xlab("Disease") +
  ylab("Cholesterol") +
  theme_fivethirtyeight() +
  theme(axis.title = element_text(), legend.position = "none")
```




Fasting blood sugar

Normal **blood sugar levels** of **non-diabetic** people are between **72mg/dL** and **99mg/dL** when fasting. Fasting blood sugar levels of **100mg/dL** up to **125mg/dL** are already described as **prediabetic**, while **fbs > 125mg/dL** are diagnosed as **diabetic**.

The study shows two possible outcomes of **fbs**: $\leq 120\text{mg/dL}$ and $>120\text{mg/dL}$. It must be considered, that people with a **fbs >120mg/dL** are at greater risk of developing heart disease or **cardiovascular disease**, however the symptoms of the patient may be caused by diabetes and secondary diseases.

Only a few patients have blood sugar levels in the range where diabetes would be diagnosed. The number of patients with and without disease are similar. The most patients have fasting blood sugar levels of 120 and lower.



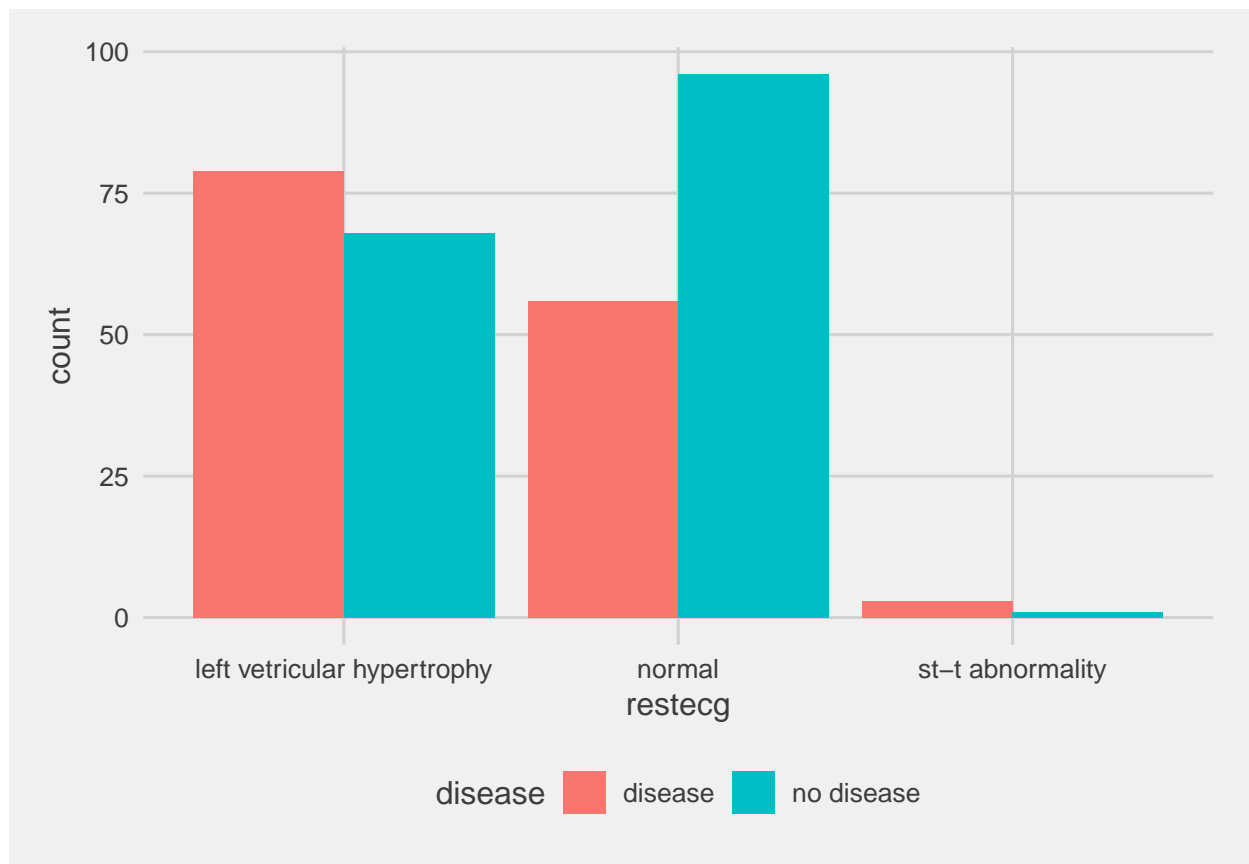
There were 14.90% of patients with a critical value of fasting blood sugar level in the database, while the prevalence of diabetes in the US was 4.90%² in the year 1990. So we can observe a much higher prevalence of patients with diabetes in the study from 1988 than in the total population.

²Diabetes trends in the U.S.: 1990-1998

Resting electrocardiographic results

As results of the **restecg** there are three potential outcomes:


- **left ventricular hypertrophy:**
Left ventricular hypertrophy is enlargement and thickening (hypertrophy) of the walls of the heart's main pumping chamber.
- **Normal:**
No abnormalities or hypertrophies.
- **Having ST-T wave abnormality:**
Abnormalities of **ST-** and/or **T wave** in the imaging procedures of the electrocardiogram.



THALACH


THALACH is the **maximum heart rate** that has been achieved of each patient.

We observe a lower maximum heart rate for patients with disease than without disease:



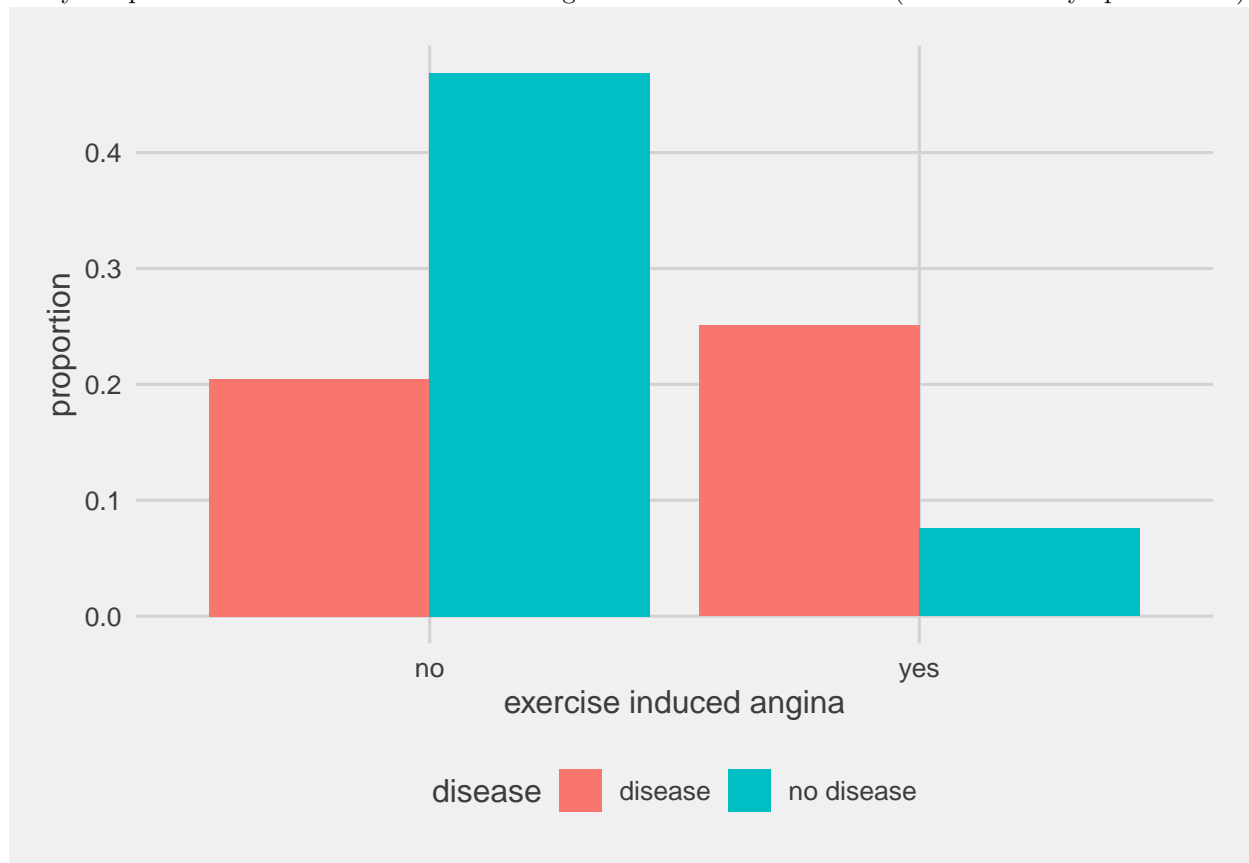
disease	mean	median
disease	139	142
no disease	158	161

As you can see in the next chart the average maximum heart rate decreases with age. An interesting abnormality is that patients with heart disease show a lower maximum heart rate at almost any age.



Exercise induced angina


We can assume that angina indicates heart disease at exercise more often than without. We can tell by the fact that most patients with exercise induced angina had heart disease but only a minority of patients with heart disease had angina outside the exercises (most were asymptomatic³).



³ chest pain type and disease

ST depression induced by exercise relative to rest

We can see that a greater **ST depression** is a sign of an increased probability of heart disease. The following findings from the database show, that the ST depression increase at exercise for patients with heart disease is greater than for patients without heart disease:

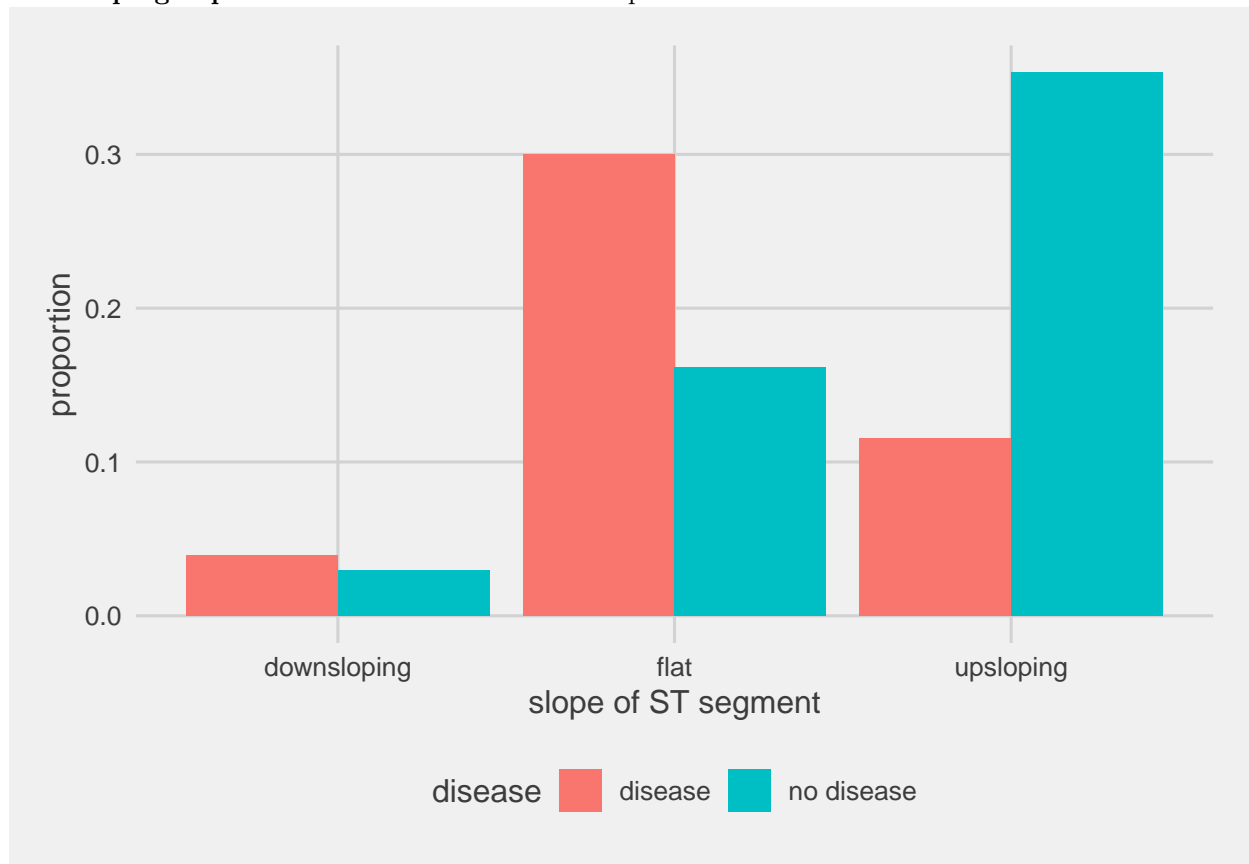


disease	mean	median
disease	1.586	1.4
no disease	0.583	0.2

Mean and median show higher values of ST depression relation in people with heart disease than in people without heart disease.


Slope of peak exercise ST segment

By the slope of **ST segment** we can see a higher proportion of patients in categories **flat slope** and **downsloping depression** that have disease than of patients without disease.




Major vessels colored by flouroscopy

Flouroscopy is an imaging tool which is made for looking on several body systems. In this case the flouroscopy was used to observe the flow of blood through three major vessels in order to evaluate the presence of arterial blockages.

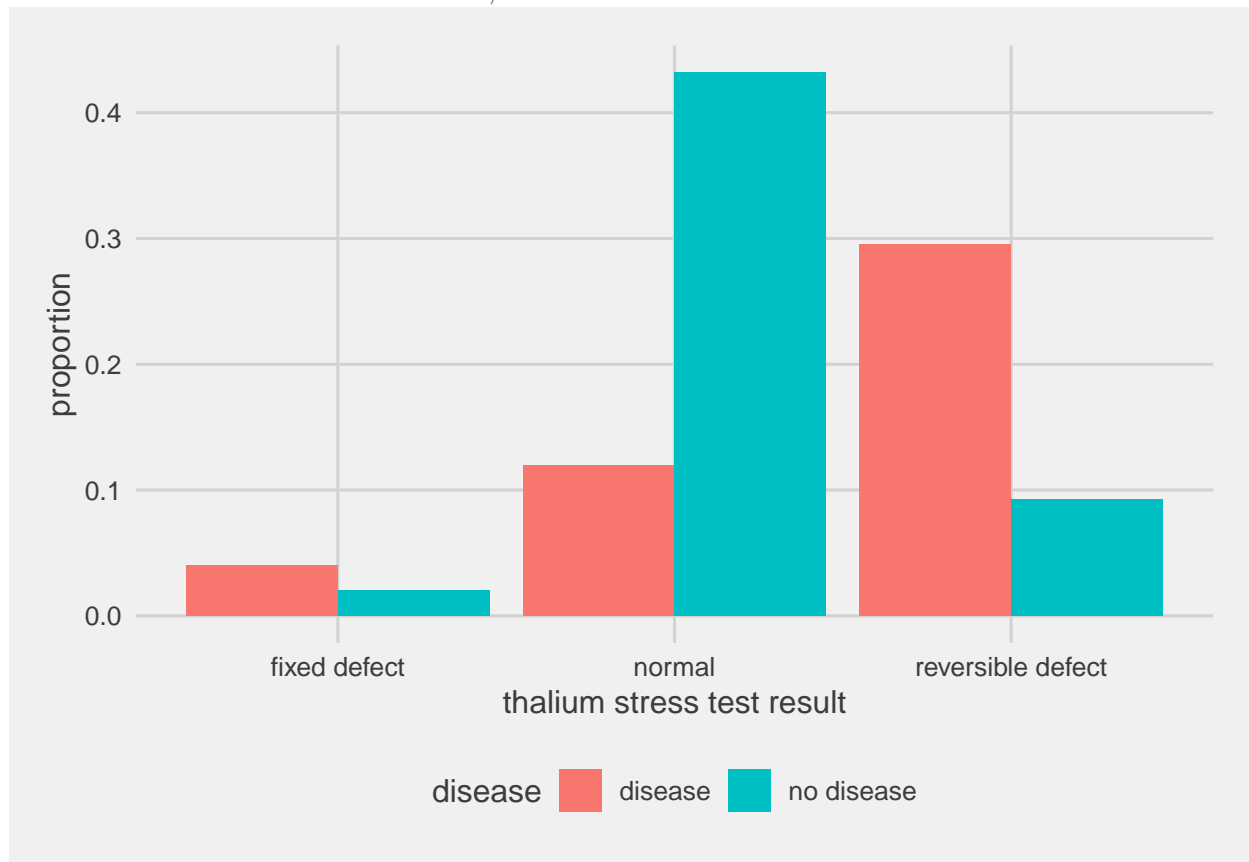


The next plot shows that the more vessels are colored by flouroscopy the higher the proportion of patients with disease.

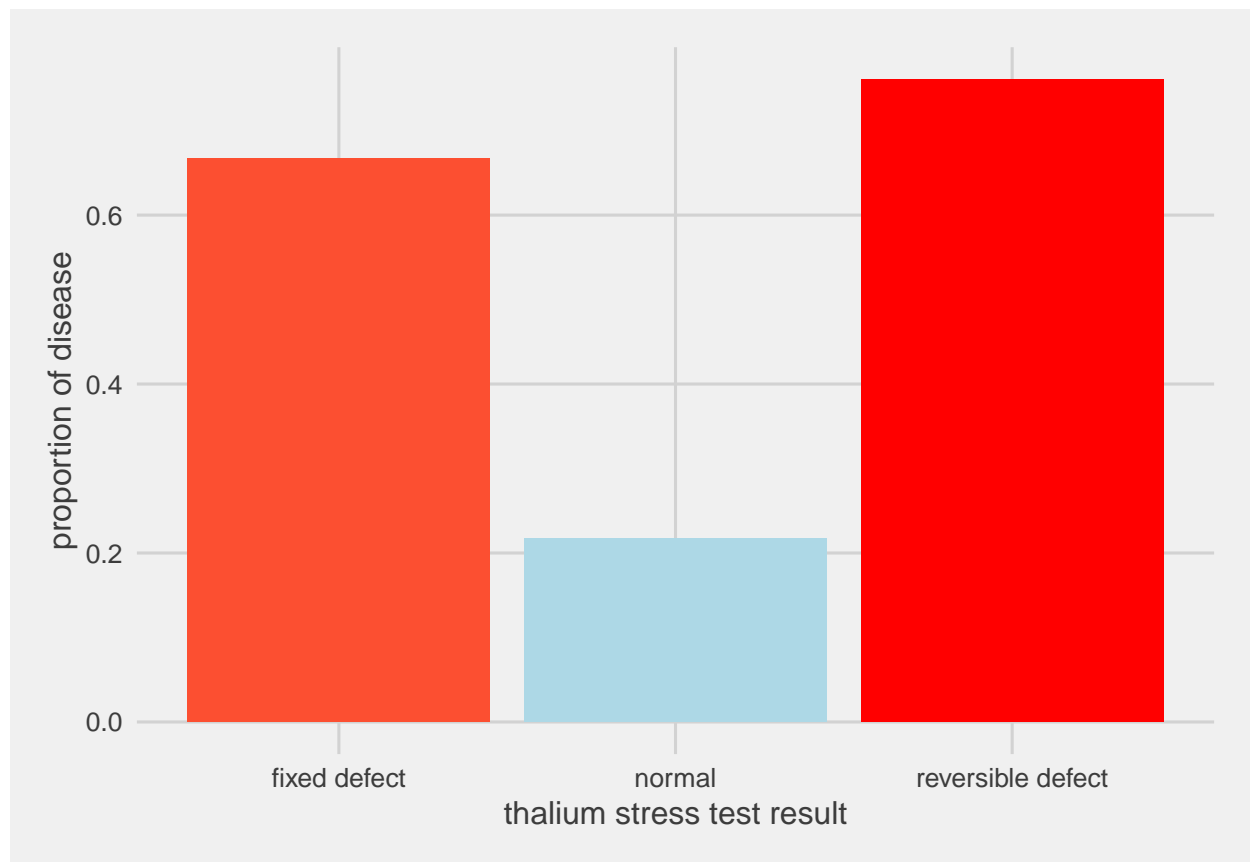


Thalium Stress Test Result

Thalium Stress Test is a test used to measure how the blood flow works while exercising or resting. The result was divided into normal result, fixed defect and reversible defect.



In the next plot we can see that the proportion of diseased patients at normal thalium stress test results is near 20%. While for any type of defect it is higher than 60%.



Methods

Training and testing set

In order to run different machine learning methods on the data we will first check the dataset on NA values. And exclude these entries from the upcoming modelings

```
sum(is.na(HeartData))
```

```
## [1] 7
```

```
compl <- as.vector(complete.cases(HeartData))  
HeartDataRM <- HeartData[compl, ]
```

To get reproducible results a seed has been set with `as.integer(Sys.time())` using the last five characters.

```
set.seed(50866, sample.kind = "default")
```

The data will be partitioned into two sets of 70% of the data for training and 30% of the data for testing.

```
test_index <- createDataPartition(y = HeartDataRM$disease,  
                                times = 1,  
                                p = 0.30,  
                                list = FALSE)  
training <- HeartDataRM[-test_index,]  
testing <- HeartDataRM[test_index,]
```

Running functions of the caret package for the next model prediction attempts already brings cross validation as default but we will use repeated cross validation with 3 repeats.

```
control.repeat <- trainControl(method = "repeatedcv",  
                              number = 10,  
                              repeats = 3)
```

To run the k-nearest neighbors algorithm later, the categorical data of the training and testing data will be encoded by the onehot encoding method using dummy variables. We will set this data to binary variables and get a data frames of 29 columns each. This data will be used for the kNN-Algorithm exclusively.

```
#one hot encoding (Training data)  
Training.dummy <- dummyVars(" ~.", data=training)  
training.onehot <- data.frame(predict(Training.dummy, newdata = training))  
training.onehot$disease.disease <- as.factor(training.onehot$disease.disease)  
training.onehot$disease.no.disease <- as.factor(training.onehot$disease.no.disease)  
training.onehot$disease.disease <- NULL  
#one hot encoding (Testing data)  
Testing.dummy <- dummyVars(" ~.", data=testing)  
testing.onehot <- data.frame(predict(Testing.dummy, newdata = testing))  
testing.onehot$disease.disease <- as.factor(testing.onehot$disease.disease)  
testing.onehot$disease.no.disease <- as.factor(testing.onehot$disease.no.disease)  
testing.onehot$disease.disease <- NULL
```

Logistic regression

Because the heart disease dataset is a categorical problem we choose the logistic regression as the first modeling approach. **Binomial** as the ‘family’-parameter indicates that the generalized linear model method is logistic regression.

```
#Logistic regression (generalized linear model)
Train.glm <- train(disease ~ ., data=training,
                  method="glm",
                  trControl=control.repeat,
                  family="binomial")
#apply model on testing
Model.glm <- predict(Train.glm, testing)
Conf.glm <- confusionMatrix(Model.glm, testing$disease)
Conf.glm$table
```

```
##           Reference
## Prediction  disease no disease
##   disease      32         5
##  no disease     9        43
```

```
Sens.glm <- Conf.glm$byClass[c("Sensitivity")]
Spec.glm <- Conf.glm$byClass[c("Specificity")]
Acc.glm <- Conf.glm$overall[["Accuracy"]]
F1.glm <- F_meas(Model.glm, testing$disease)
Prec.glm <- Conf.glm$byClass[c("Precision")]
Prev.glm <- Conf.glm$byClass[c("Prevalence")]
```

Decision Tree

Decision trees are pretty suitable for the purpose of identifying if several indicators implicate diseases or not. We have used the train function from the caret package and set the rpart method. TuneLength is set to 10, which means that the function uses ten different hyperparameters and chooses the best fitting for the training data. The hyperparameter of rpart is the **complexity parameter**.

```
Train.dec.tree <- train(disease ~ ., data=training,
  method="rpart",
  trControl=control.repeat,
  tuneLength=10
)
Model.dec.tree <- predict(Train.dec.tree, testing, type="raw")
Conf.dec.tree <- confusionMatrix(table(Model.dec.tree, testing$disease))
```

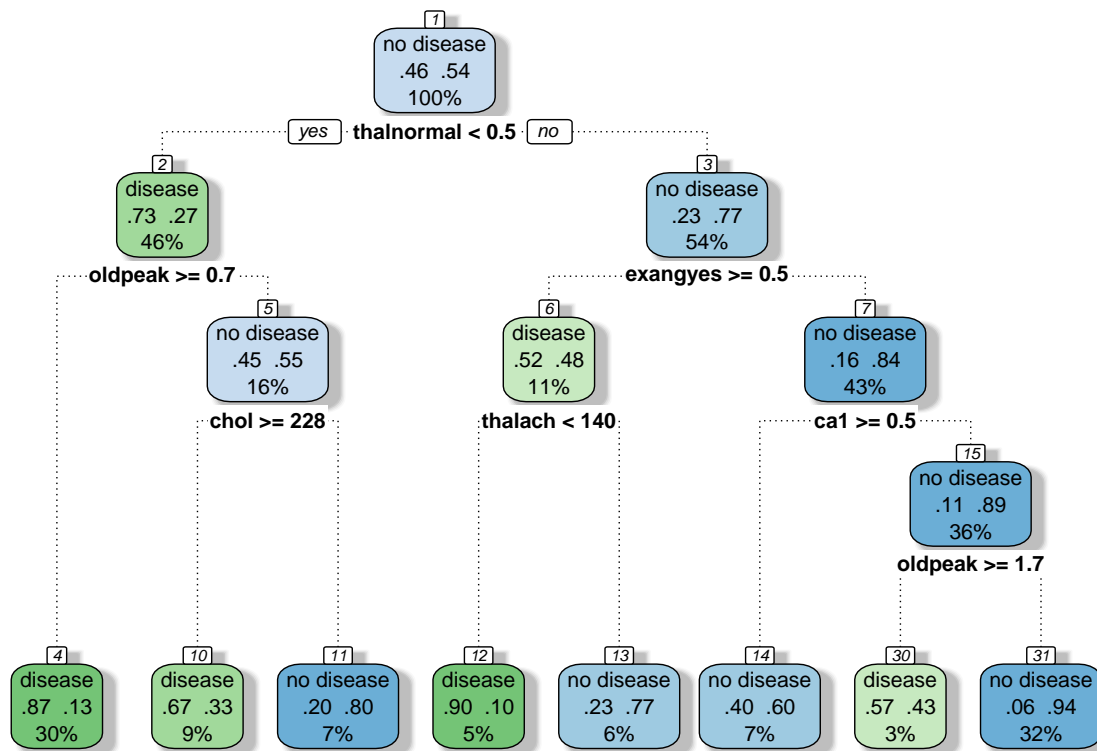
- 82.90% of patients would have been diagnosed correctly to have heart disease (sensitivity)
- 77.10% of patients would have been diagnosed correctly to have no heart disease (specificity)
- The overall accuracy of the decision tree is 79.80%.

	disease	no disease
disease	34	11
no disease	7	37

Method	Sensitivity	Specificity	Accuracy
Decision tree	0.829	0.771	0.798

The following classification tree shows that the data was first split at the result of **thal**. Furthermore split at **oldpeak**, **chol**, **exang**, **thalach** and **ca**.

- 30% of all patients were predicted to have an abnormal thal and an oldpeak of ≥ 0.7
 - 87% of these patients were predicted to have heart disease
- 9% of patients were predicted to have an abnormal thal, an oldpeak of ≥ 0.7 and chol ≥ 228
 - 67% of these patients were predicted to have heart disease
- 6% of all patients were predicted to have a normal thal, exercise induced angina and thalach ≥ 140
 - 77% of these patients were predicted to have no heart disease
- 32% of the patients were predicted to have a normal thal, exercise induced angina, ca $\neq 1$ and have oldpeak ≥ 1.7
 - 94% of these patients were predicted to have no heart disease



*Split of categorical data is between 1 (true) and 0 (false) - e.g. thalnormal < 0.5 means all abnormal thal results.

Random forest

With the random forest as an extension of the decision tree there is a high potential in outperforming a single tree by using several hundred trees. The number of trees is set to default (n=500).

```
Train.random.forest <- train(disease ~ ., data=training,
                             method="rf",
                             metric="Accuracy",
                             preProcess=c("center", "scale"),
                             tuneLength=10,
                             trControl=control.repeat
                             )
Model.random.forest <- predict(Train.random.forest, testing, type="raw")
Conf.random.forest <- confusionMatrix(Model.random.forest, testing$disease)
```

	disease	no disease
disease	34	8
no disease	7	40


Method	Sensitivity	Specificity	Accuracy
Random forest	0.829	0.833	0.831

As we can see the random forest outperforms the decision tree. Now we run this function with different amounts of ntree:

```
modellist <- list()
for (ntree in c(500, 1000, 1500, 2000, 2500)) {
  fit <- train(disease ~ ., data=training,
              method="rf",
              metric="Accuracy",
              preProcess=c("center", "scale"),
              tuneGrid=expand.grid(.mtry=c(sqrt(ncol(training)))),
              trControl=control.repeat,
              ntree=ntree)
  key <- toString(ntree)
  modellist[[key]] <- fit
}
results <- resamples(modellist)
```

In the next plot we can see that the highest accuracy of the random forest is located at an amount of 2500 trees.

```
dotplot(results, ylab="nTree")
```



So we correct the amount of trees upwards.

```
Train.random.forest <- train(disease ~ ., data=training,
                             method="rf",
                             metric="Accuracy",
                             preProcess=c("center", "scale"),
                             ntree=2500,
                             tuneLength=10,
                             trControl=control.repeat
                             )
Model.random.forest <- predict(Train.random.forest, testing, type="raw")
Conf.random.forest <- confusionMatrix(Model.random.forest, testing$disease)
```

	disease	no disease
disease	34	9
no disease	7	39

Method	Sensitivity	Specificity	Accuracy
Random forest	0.829	0.812	0.82

Finally for random forest we have a look on the importance score of each variable in the dataset. We have already seen the attributes with the highest importance in the decision tree as decision at quite high nodes.

```
## rf variable importance
##
```

##	Overall
## oldpeak	100.00
## thalach	77.24
## age	60.67
## thalnormal	59.53
## chol	58.26
## thalreversible defect	53.05
## trestbps	50.85
## ca1	36.23
## slopeupsloping	32.87
## exangyes	31.90
## sexmale	27.97
## cpnon anginal pain	27.26
## ca2	22.86
## slopeflat	20.79
## restecgnormal	16.18
## cptypical angina	12.82
## ca3	9.26
## cpatypical angina	8.69
## fbs>120	5.09
## restecgst-t abnormality	0.00

Support vector machines

With the support vector machines we choose a linear method. Using a tuneLength of 5 the train function searches for the optimal tuning parameter of cost.

```
train.svmLinear <- train(disease ~ ., data=training,
  method = "svmLinear",
  trControl= control.repeat,
  preProcess=c("center",
    "scale"),
  # tuneGrid=expand.grid(C=1)
  tuneLength=5)
Model.svmLinear <- predict(train.svmLinear, testing, type="raw")
Conf.svmLinear <- confusionMatrix(Model.svmLinear, testing$disease)
```

	disease	no disease
disease	27	7
no disease	14	41

Method	Sensitivity	Specificity	Accuracy
Svm linear	0.659	0.854	0.764

K-nearest neighbors

Another modeling approach in this project is the k-nearest neighbors model. Using a tuneLength of 3 the train function searches for the optimal size of neighbors to be based on as a tuning parameter.

```
Train.knn <- train(disease.no.disease~., data=training.onehot,  
  method="knn",  
  trControl=control.repeat,  
  # tuneGrid=expand.grid(k=5)  
  tuneLength=3  
)  
Model.knn <- predict(Train.knn, testing.onehot, type = "raw")  
Conf.knn <- confusionMatrix(Model.knn, testing.onehot$disease)
```

	0	1
0	24	8
1	17	40

- disease (0); no disease (1)

Method	Sensitivity	Specificity	Accuracy
K-nearest neighbors	0.585	0.833	0.719

Results/Evaluation metrics

With different models passed through there is now a bunch of metrics by which we can justify what model is the best to use for the heart disease problem. Besides collecting standard measures for evaluating different models for the problem, we have to take a closer look on the problem itself and its sector of use and consequences of using the model in order to treat patients right.

Method	Sensitivity/Recall	Specificity	Accuracy	F1-score	Precision
Logistic regression	0.780	0.896	0.843	0.821	0.865
Decision tree	0.829	0.771	0.798	0.791	0.756
Random forest	0.829	0.812	0.820	0.810	0.791
Support vector machine (linear)	0.659	0.854	0.764	0.720	0.794
k-nearest neighbors	0.585	0.833	0.719	0.658	0.750

Usually the focus of prediction approaches in the medical field should be on **reducing false negative results** to prevent overlooking diseases of a patient. Therefore receiving false negative cases is worse than receiving false positive cases. With this approach we will assess the different models.

Sensitivity/Recall also called true positive rate provides information about how precise the model is in order of finding all patients with heart disease. We can see that the decision tree as well as the random forest have the highest sensitivity of 82.9% which means, that eight out of ten patients that have heart disease were diagnosed correctly. K-nearest neighbors estimate shows a poor result for sensitivity, only diagnosing six out of ten patients correctly. In terms of specificity, all methods show a good result.

Logistic regression and random forest are the only methods that show an overall accuracy over 80.0% as well as a F1-score over 80.0%.

With a focus on keeping false negative cases as low as possible the random forest brings solid measurement accuracies over 80.0% in every single metric.

Conclusion

Since some methods like random forest and logistic regression show sufficient results, the attributes of the dataset seem to be good indicators to predict heart disease. However the dataset was released in 1988 around 40 years ago. It can be assumed that there are additional attributes to diagnose patients to have heart disease but more precisely.